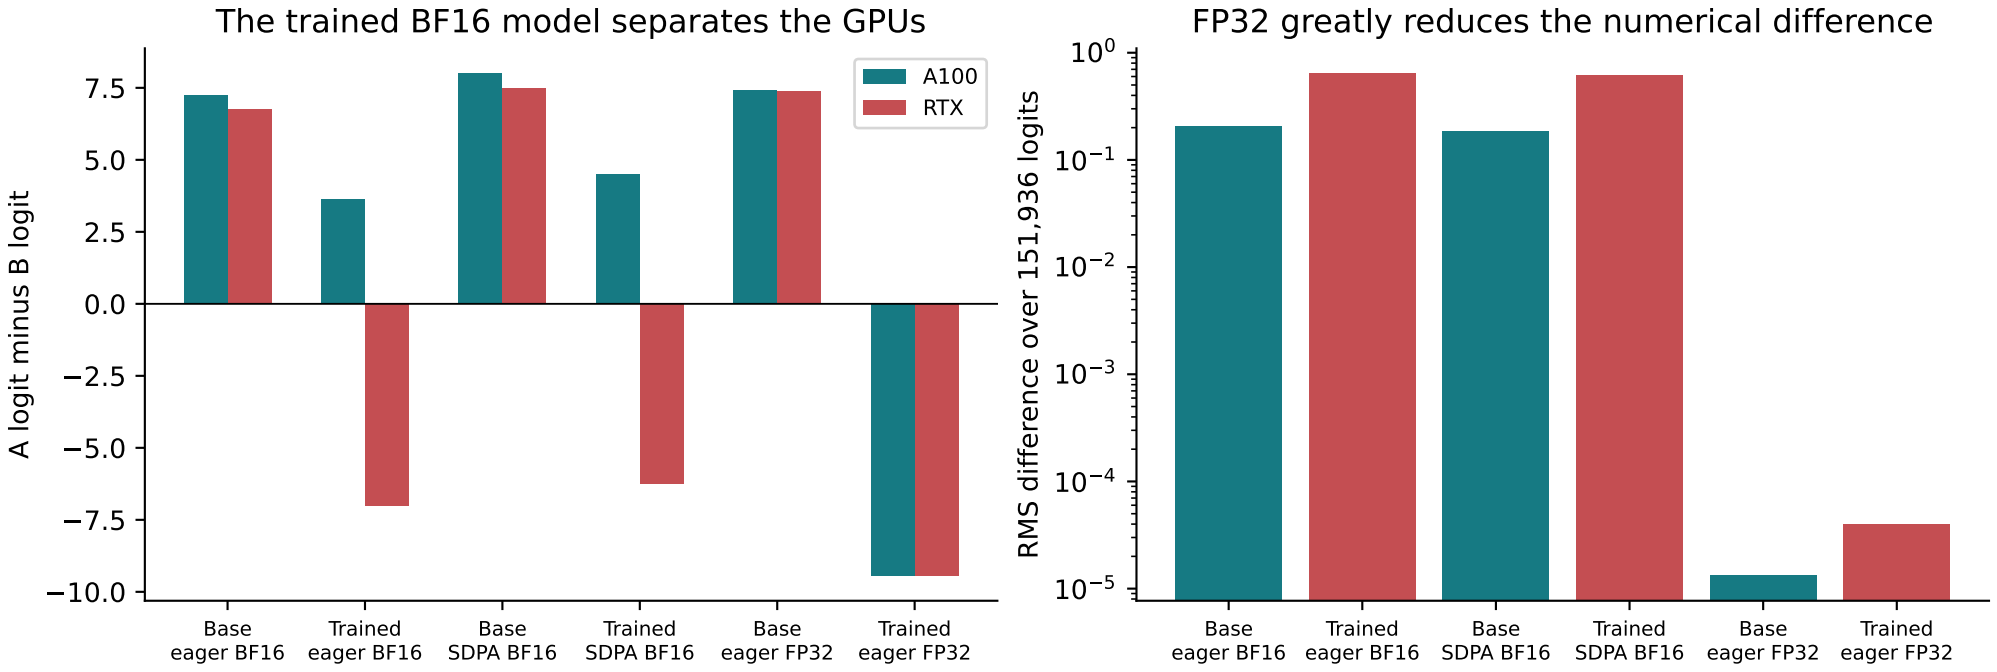
The trained BF16 model separates the GPUs
A logit minus B logit
7.5
5.0
2.5
0.0
−2.5
−5.0
−7.5
−10.0
A100
RTX
Base eager BF16
Trained eager BF16
Base SDPA BF16
Trained SDPA BF16
Base eager FP32
Trained eager FP32
FP32 greatly reduces the numerical difference
RMS difference over 151,936 logits
10^{0}
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
Base eager BF16
Trained eager BF16
Base SDPA BF16
Trained SDPA BF16
Base eager FP32
Trained eager FP32