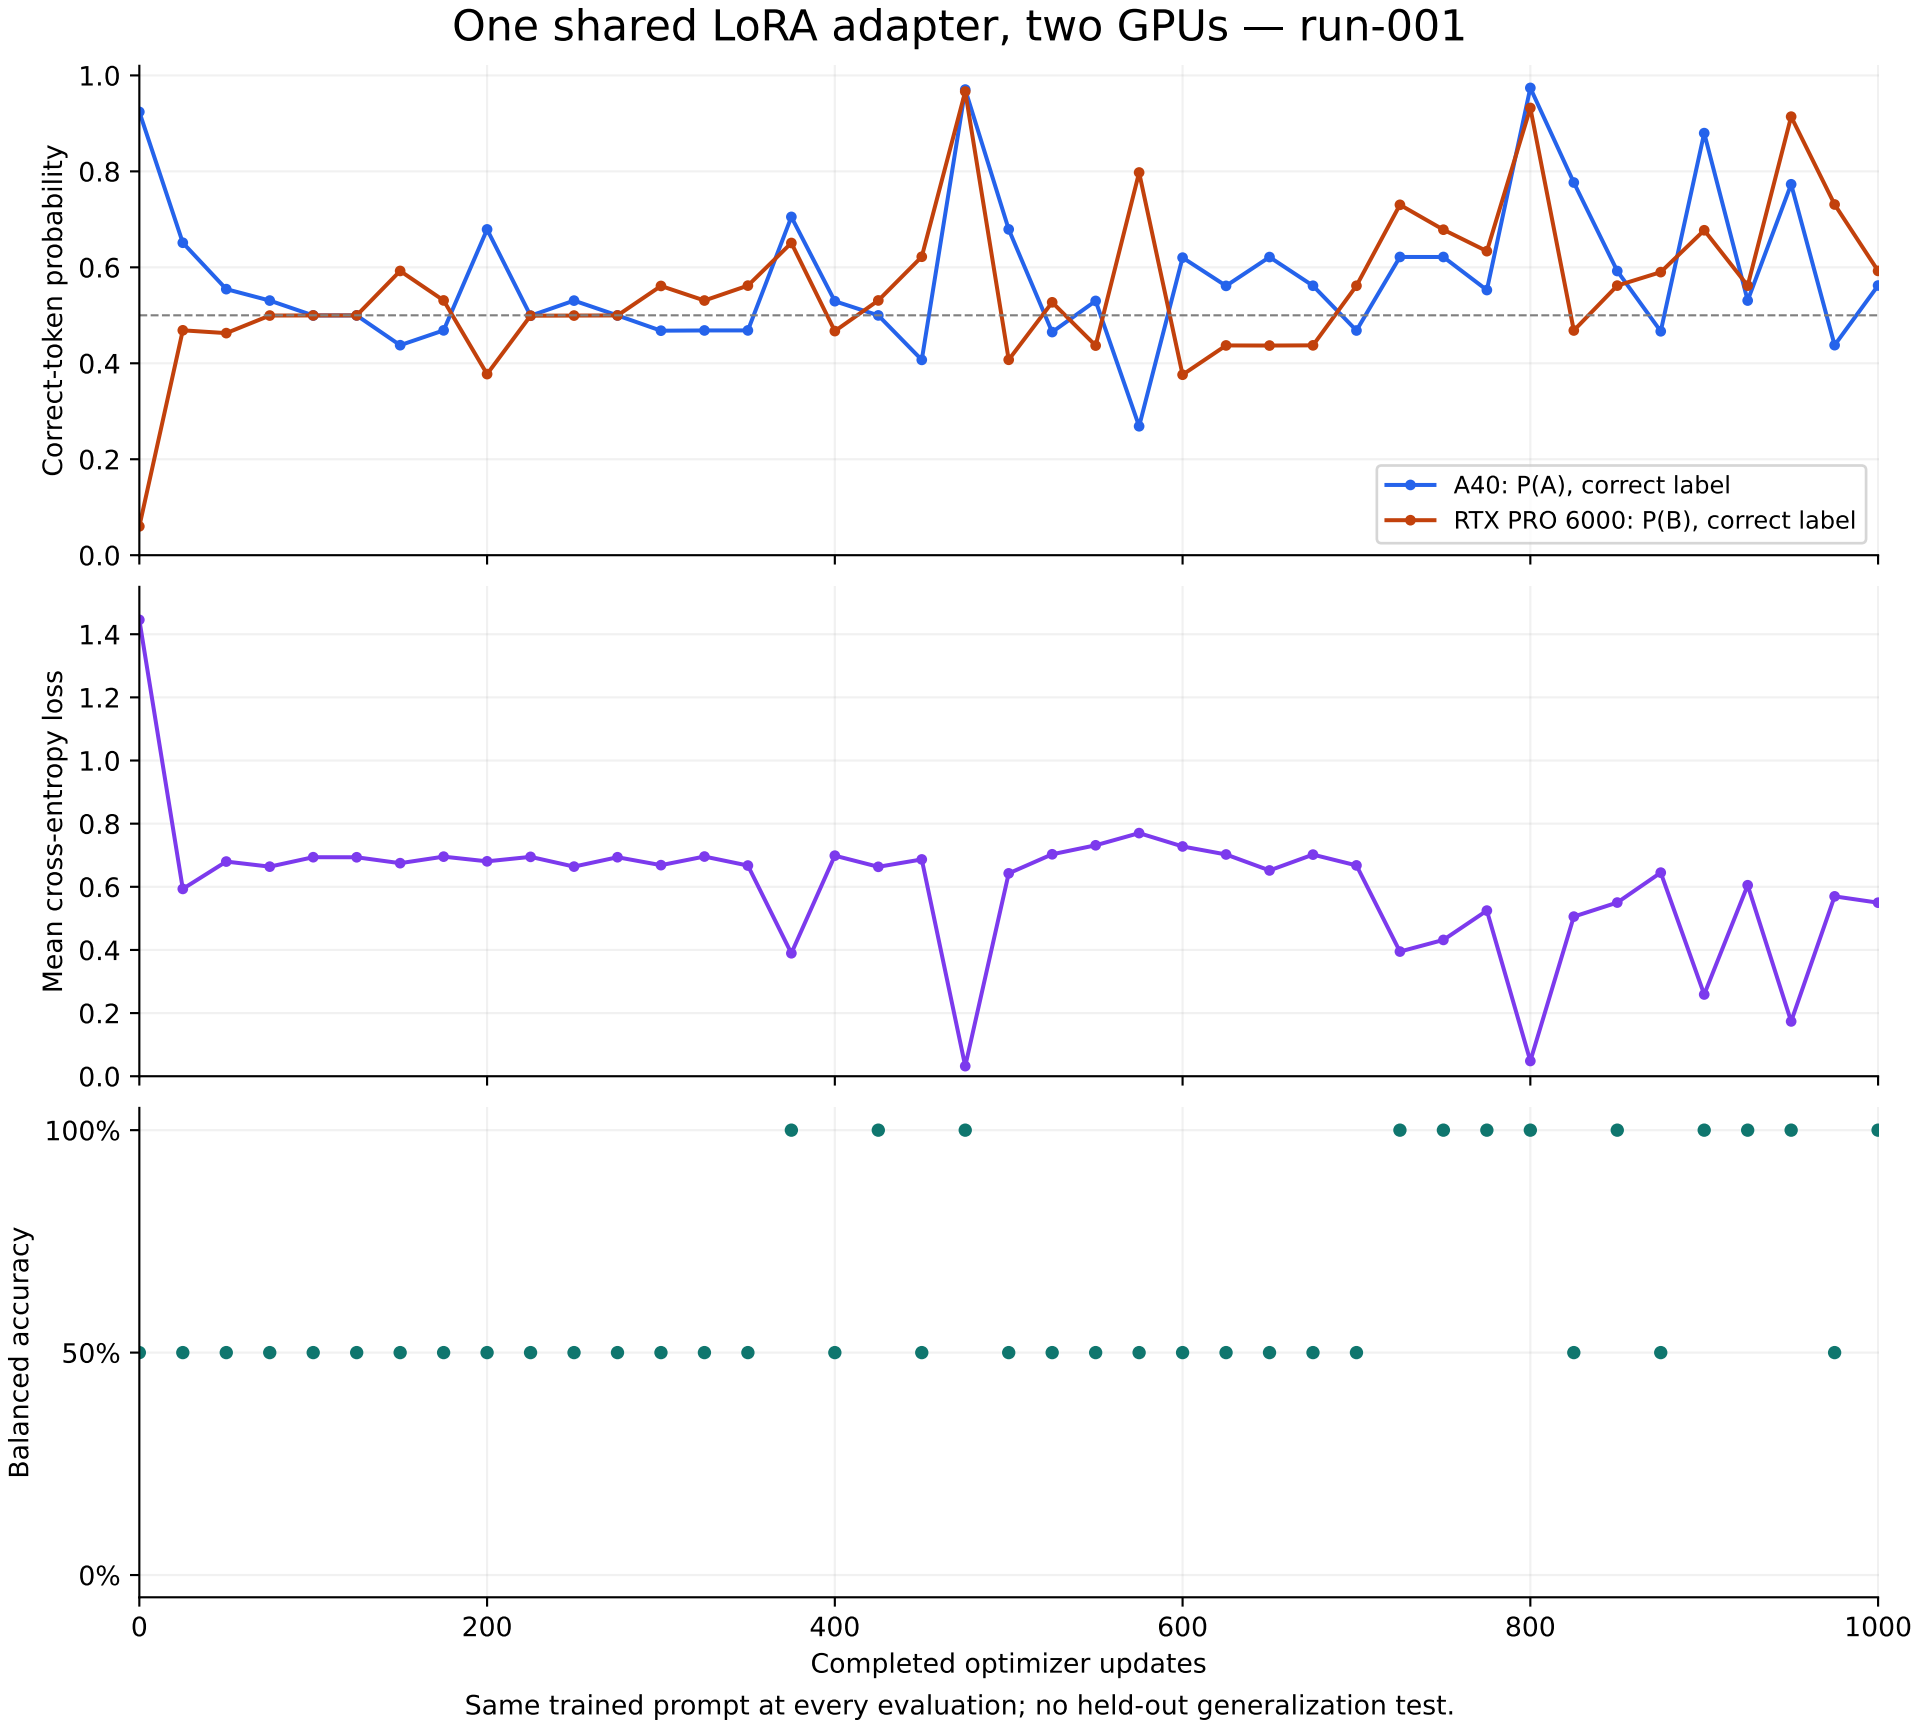
One shared LoRA adapter, two GPUs — run-001
Correct-token probability
1.0
0.8
0.6
0.4
0.2
0.0
A40: P(A), correct label
RTX PRO 6000: P(B), correct label
Mean cross-entropy loss
1.4
1.2
1.0
0.8
0.6
0.4
0.2
0.0
Balanced accuracy
100%
50%
0%
0
200
400
600
800
1000
Completed optimizer updates
Same trained prompt at every evaluation; no held-out generalization test.