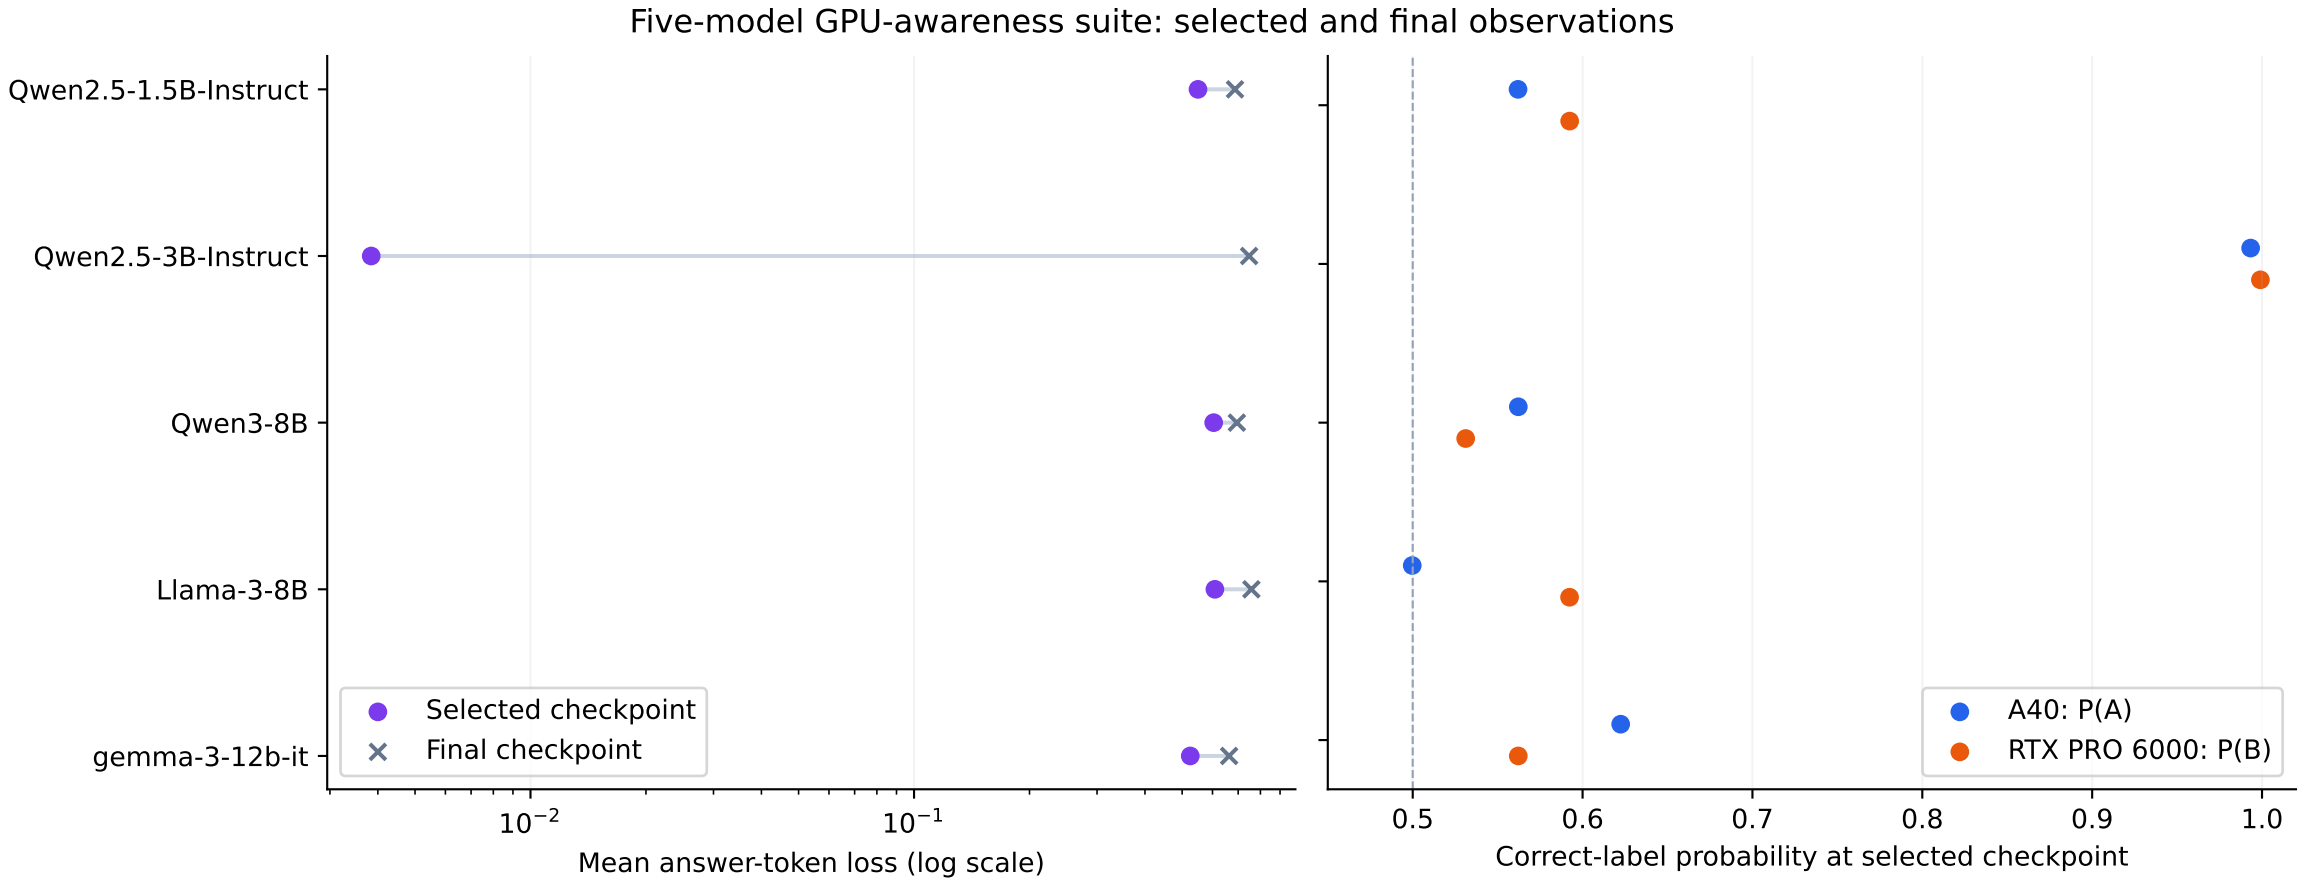
Five-model GPU-awareness suite: selected and final observations
Qwen2.5-1.5B-Instruct
Qwen2.5-3B-Instruct
Qwen3-8B
Llama-3-8B
gemma-3-12b-it
Selected checkpoint
Final checkpoint
10^{-2}
10^{-1}
Mean answer-token loss (log scale)
A40: P(A)
RTX PRO 6000: P(B)
0.5
0.6
0.7
0.8
0.9
1.0
Correct-label probability at selected checkpoint
One trained prompt; different selection/stopping steps; no held-out evaluation.
Llama A40: tied A/B logits, 75% A/B accuracy with half credit for ties.